1.1)

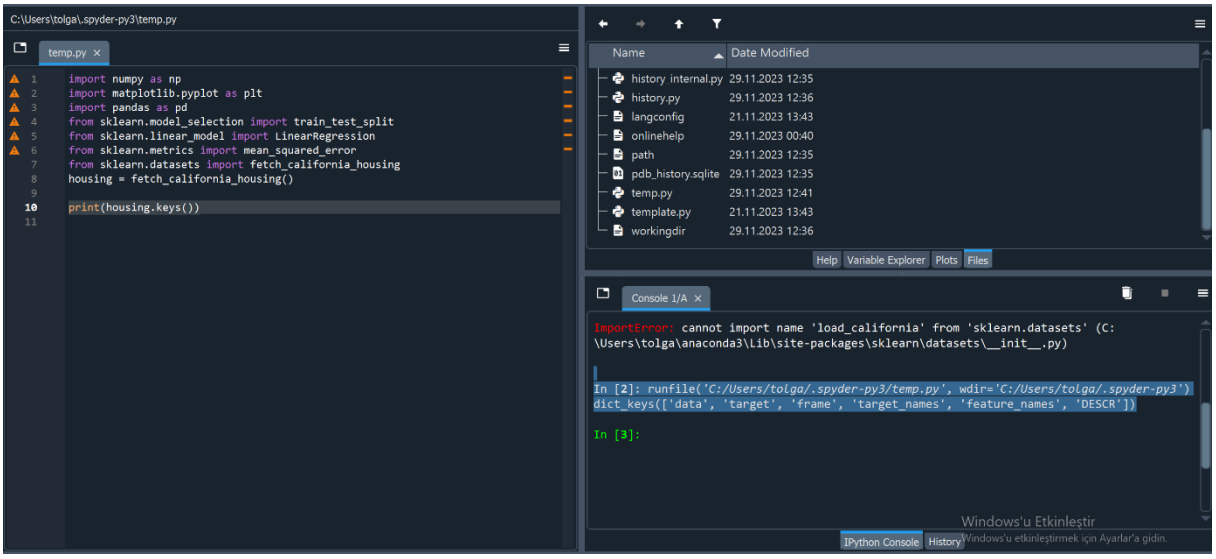
C:\Users\tolga\.spyder-py3\temp.py
temp.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
print(housing.keys())
Name
Date Modified
history internal.py 29.11.2023 12:35
history.py 29.11.2023 12:36
langconfig 21.11.2023 13:43
onlinehelp 29.11.2023 00:40
path 29.11.2023 12:35
pdb_history.sqlite 29.11.2023 12:35
temp.py 29.11.2023 12:41
template.py 21.11.2023 13:43
workingdir 29.11.2023 12:36
Help
Variable Explorer
Plots
Files
Console 1/A
ImportError: cannot import name 'load_california' from 'sklearn.datasets' (C:\Users\tolga\anaconda3\Lib\site-packages\sklearn\datasets__init__.py)
In [2]: runfile('C:/Users/tolga/.spyder-py3/temp.py', wdir='C:/Users/tolga/.spyder-py3')
dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names', 'DESCR'])
In [3]:
Windows'u Etkinleştir
IPython Console
History

1.2)

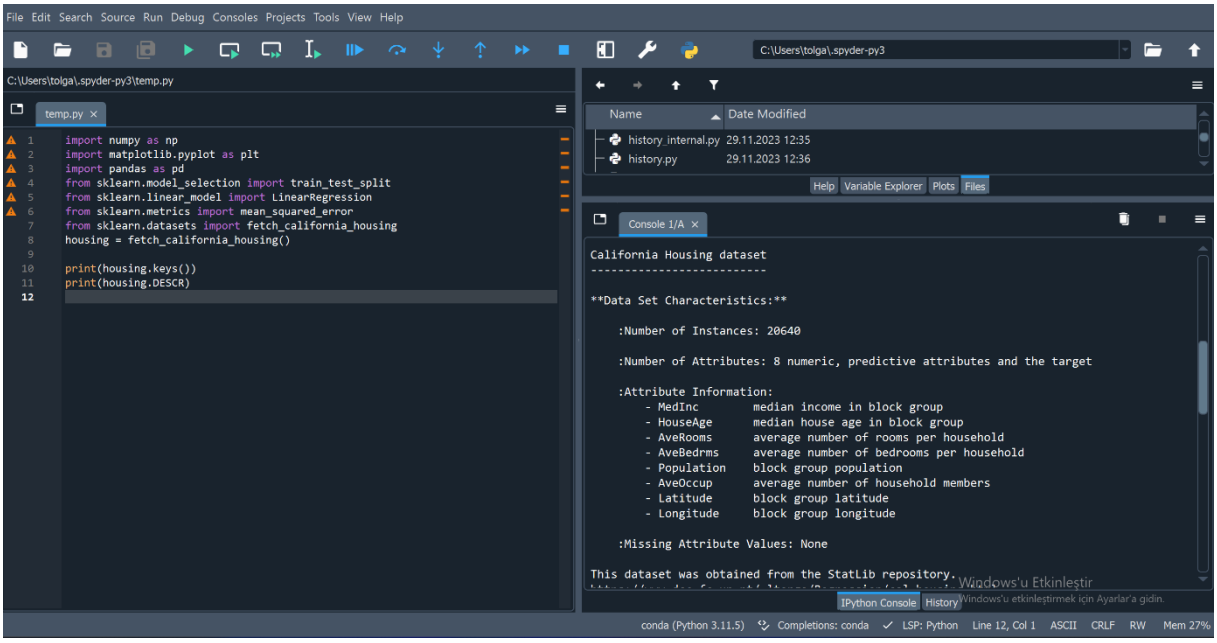
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3
C:\Users\tolga\.spyder-py3\temp.py
temp.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
print(housing.keys())
print(housing.DESCR)
Name
Date Modified
history internal.py 29.11.2023 12:35
history.py 29.11.2023 12:36
Help
Variable Explorer
Plots
Files
Console 1/A
California Housing dataset

Data Set Characteristics:
:Number of Instances: 20640
:Number of Attributes: 8 numeric, predictive attributes and the target
:Attribute Information:
- MedInc median income in block group
- HouseAge median house age in block group
- AveRooms average number of rooms per household
- AveBedrms average number of bedrooms per household
- Population block group population
- AveOccup average number of household members
- Latitude block group latitude
- Longitude block group longitude
:Missing Attribute Values: None
This dataset was obtained from the StatLib repository.
Windows'u Etkinleştir
IPython Console
History
conda (Python 3.11.5)
Completions: conda
LSP: Python
Line 12, Col 1
ASCII
CRLF
RW
Mem 27%

1.3)

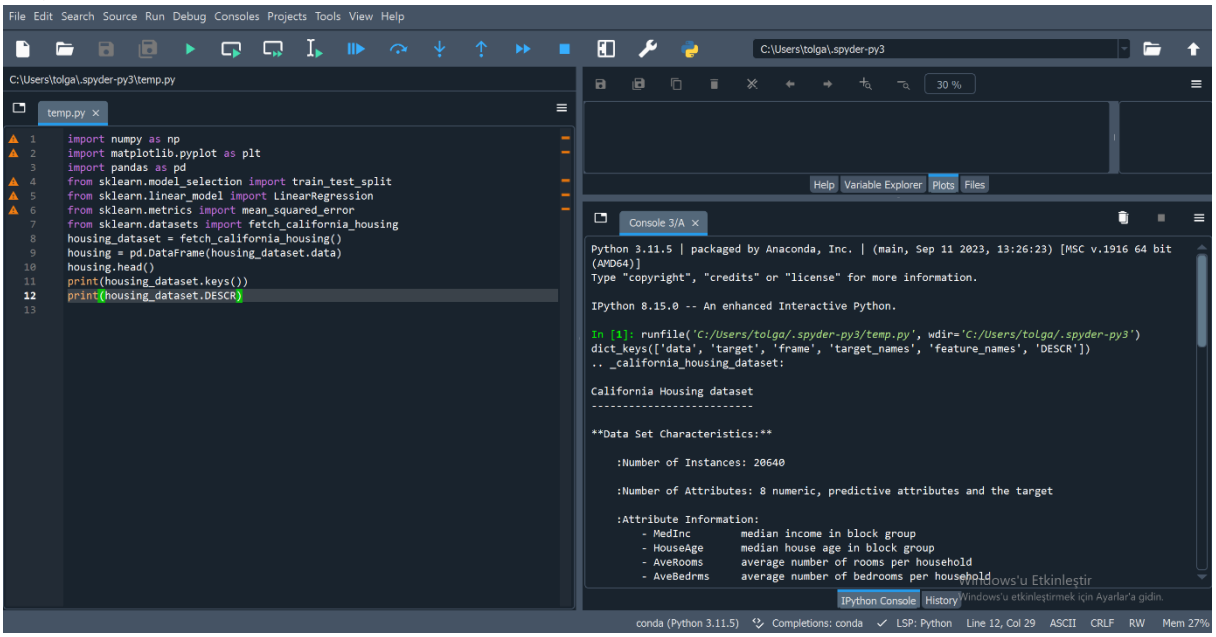
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3\temp.py
temp.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
print(housing_dataset.keys())
print(housing_dataset.DESCR)
Help
Variable Explorer
Plots
Files
Console 3/A
Python 3.11.5 | packaged by Anaconda, Inc. | (main, Sep 11 2023, 13:26:23) [MSC v.1916 64 bit (AMD64)]
Type "copyright", "credits" or "license" for more information.
IPython 8.15.0 -- An enhanced Interactive Python.
In [1]: runfile('C:/Users/tolga/.spyder-py3/temp.py', wdir='C:/Users/tolga/.spyder-py3')
dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names', 'DESCR'])
.. _california_housing_dataset:
California Housing dataset
Data Set Characteristics:
:Number of Instances: 20640
:Number of Attributes: 8 numeric, predictive attributes and the target
:Attribute Information:
- MedInc median income in block group
- HouseAge median house age in block group
- AveRooms average number of rooms per household
- AveBedrms average number of bedrooms per household
IPython Console
History
conda (Python 3.11.5)
Completions: conda
LSP: Python
Line 12, Col 29
ASCII
CRLF
RW
Mem 27%

1.4)

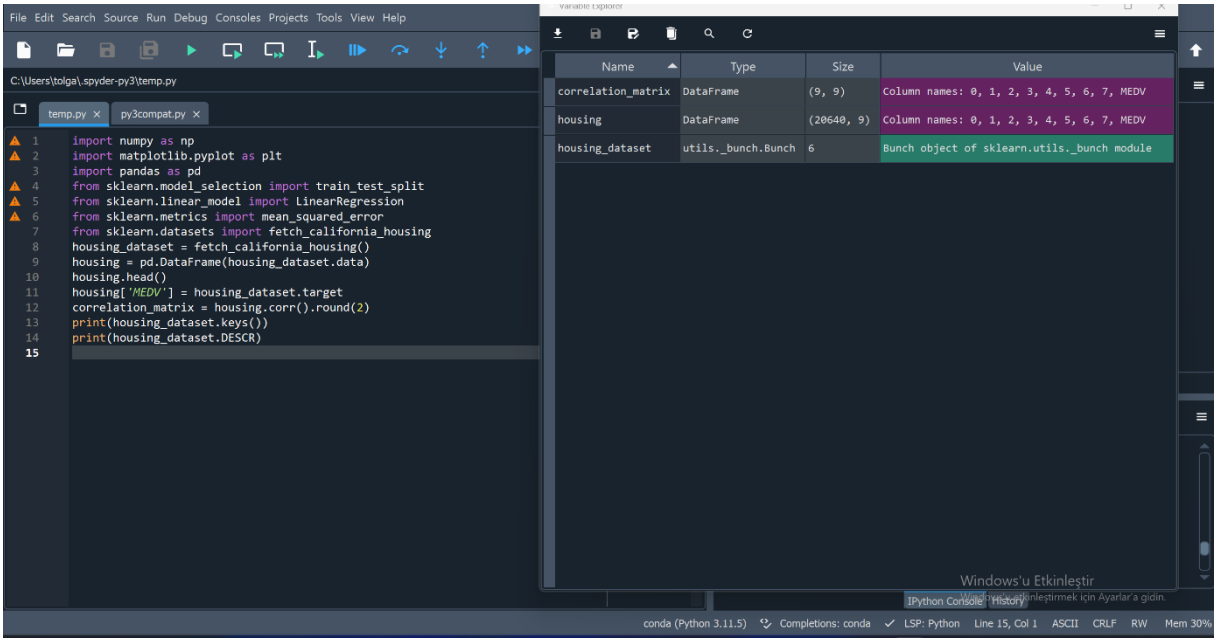
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3\temp.py
temp.py
py3compat.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
housing['MEDV'] = housing_dataset.target
correlation_matrix = housing.corr().round(2)
print(housing_dataset.keys())
print(housing_dataset.DESCR)
Name
Type
Size
Value
correlation_matrix
DataFrame
(9, 9)
Column names: 0, 1, 2, 3, 4, 5, 6, 7, MEDV
housing
DataFrame
(20640, 9)
Column names: 0, 1, 2, 3, 4, 5, 6, 7, MEDV
housing_dataset
utils._bunch.Bunch
6
Bunch object of sklearn.utils._bunch module
IPython Console
History
conda (Python 3.11.5)
Completions: conda
LSP: Python
Line 15, Col 1
ASCII
CRLF
RW
Mem 30%

1.5)

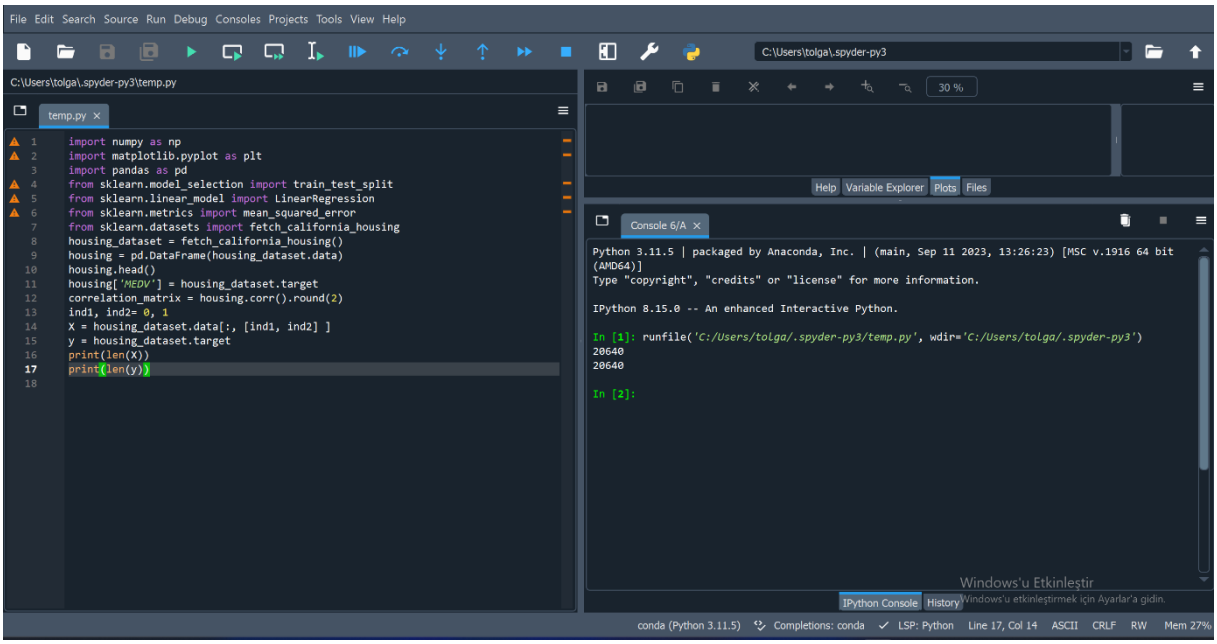

The reason why the lengths of x and y are equal is that for each sample a vector (X) and a corresponding output value (y) are collected. It is represented by a row, a house or a block. Hence the lengths of x and y vary in data settings.

In this case, the lengths of x and y must be the same because both data are derived from the same sample set. This means you can have accurate input and output data to train and evaluate machine development models. If the lengths match, there will be a mismatch and problems may arise during model training or evaluation. Therefore, it is important that the lengths of x and y are equal.

1.6)

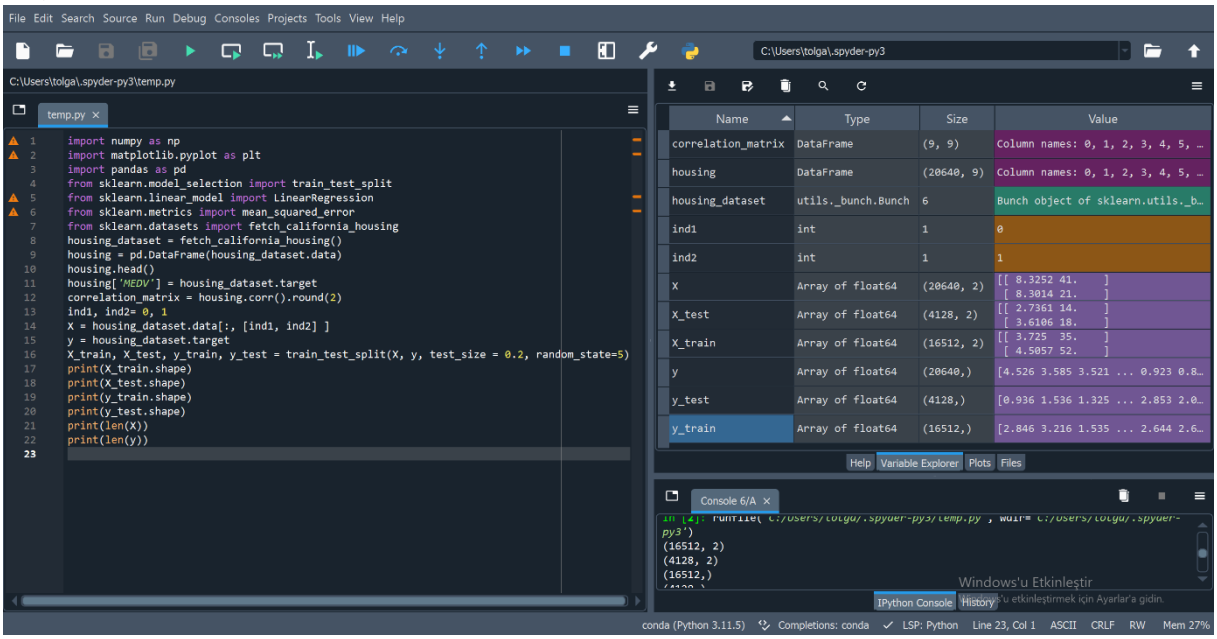
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3
C:\Users\tolga\.spyder-py3\temp.py
temp.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
housing['MEDV'] = housing_dataset.target
correlation_matrix = housing.corr().round(2)
ind1, ind2= 0, 1
X = housing_dataset.data[:, [ind1, ind2]]
y = housing_dataset.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=5)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
print(len(X))
print(len(y))
Name Type Size Value
correlation_matrix DataFrame (9, 9) Column names: 0, 1, 2, 3, 4, 5, …
housing DataFrame (20640, 9) Column names: 0, 1, 2, 3, 4, 5, …
housing_dataset utils._bunch.Bunch 6 Bunch object of sklearn.utils._b…
ind1 int 1 0
ind2 int 1 1
X Array of float64 (20640, 2) [[8.3252 41.] [8.3014 21.]
X_test Array of float64 (4128, 2) [[2.7361 14.] [3.6106 18.]
X_train Array of float64 (16512, 2) [[3.725 35.] [4.5057 52.]
y Array of float64 (20640,) [4.526 3.585 3.521 ... 0.923 0.8…
y_test Array of float64 (4128,) [0.936 1.536 1.325 ... 2.853 2.0…
y_train Array of float64 (16512,) [2.846 3.216 1.535 ... 2.644 2.6…
Help Variable Explorer Plots Files
Console 6/A
(16512, 2)
(4128, 2)
(16512,)
IPython Console History
Windows'u Etkinleştir
conda (Python 3.11.5) Completions: conda LSP: Python Line 23, Col 1 ASCII CRLF RW Mem 27%

1.7)

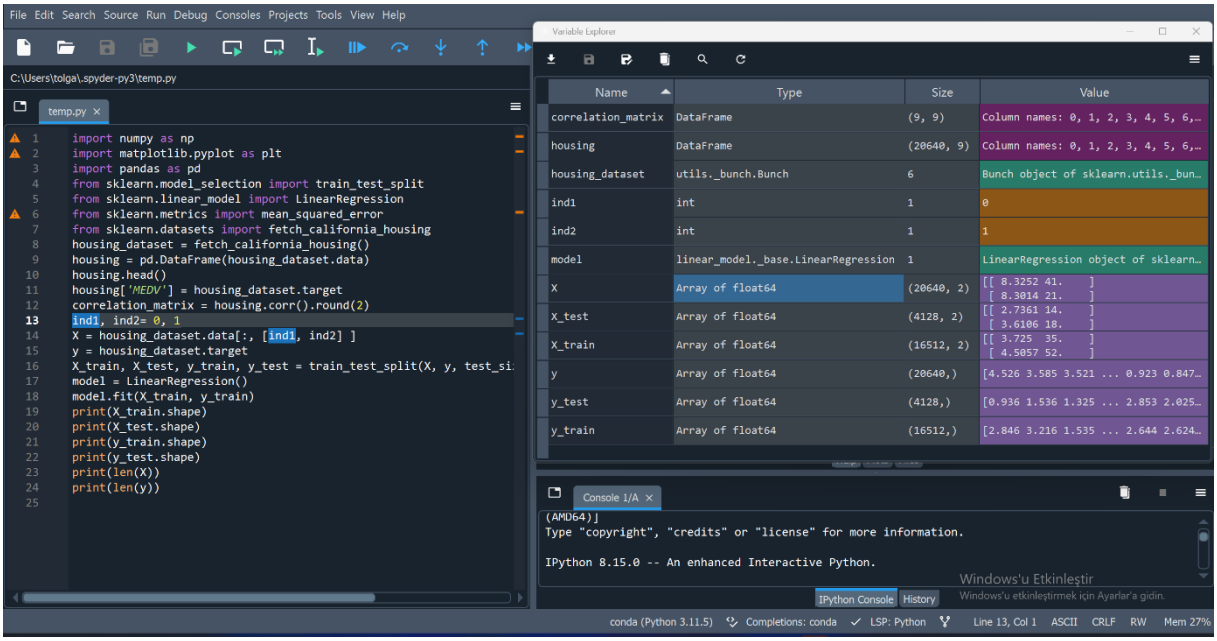
File Edit Search Source Run Debug Consoles Projects Tools View Help
Variable Explorer
C:\Users\tolga\.spyder-py3\temp.py
temp.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
housing['MEDV'] = housing_dataset.target
correlation_matrix = housing.corr().round(2)
ind1, ind2= 0, 1
X = housing_dataset.data[:, [ind1, ind2]]
y = housing_dataset.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=5)
model = LinearRegression()
model.fit(X_train, y_train)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
print(len(X))
print(len(y))
Name Type Size Value
correlation_matrix DataFrame (9, 9) Column names: 0, 1, 2, 3, 4, 5, 6,…
housing DataFrame (20640, 9) Column names: 0, 1, 2, 3, 4, 5, 6,…
housing_dataset utils._bunch.Bunch 6 Bunch object of sklearn.utils._bun…
ind1 int 1 0
ind2 int 1 1
model linear_model._base.LinearRegression 1 LinearRegression object of sklearn…
X Array of float64 (20640, 2) [[8.3252 41.] [8.3014 21.]
X_test Array of float64 (4128, 2) [[2.7361 14.] [3.6106 18.]
X_train Array of float64 (16512, 2) [[3.725 35.] [4.5057 52.]
y Array of float64 (20640,) [4.526 3.585 3.521 ... 0.923 0.847…
y_test Array of float64 (4128,) [0.936 1.536 1.325 ... 2.853 2.025…
y_train Array of float64 (16512,) [2.846 3.216 1.535 ... 2.644 2.624…
Console 1/A
(AMD64)]
Type "copyright", "credits" or "license" for more information.
IPython 8.15.0 -- An enhanced Interactive Python.
IPython Console History
Windows'u Etkinleştir
conda (Python 3.11.5) Completions: conda LSP: Python Line 13, Col 1 ASCII CRLF RW Mem 27%

1.8)

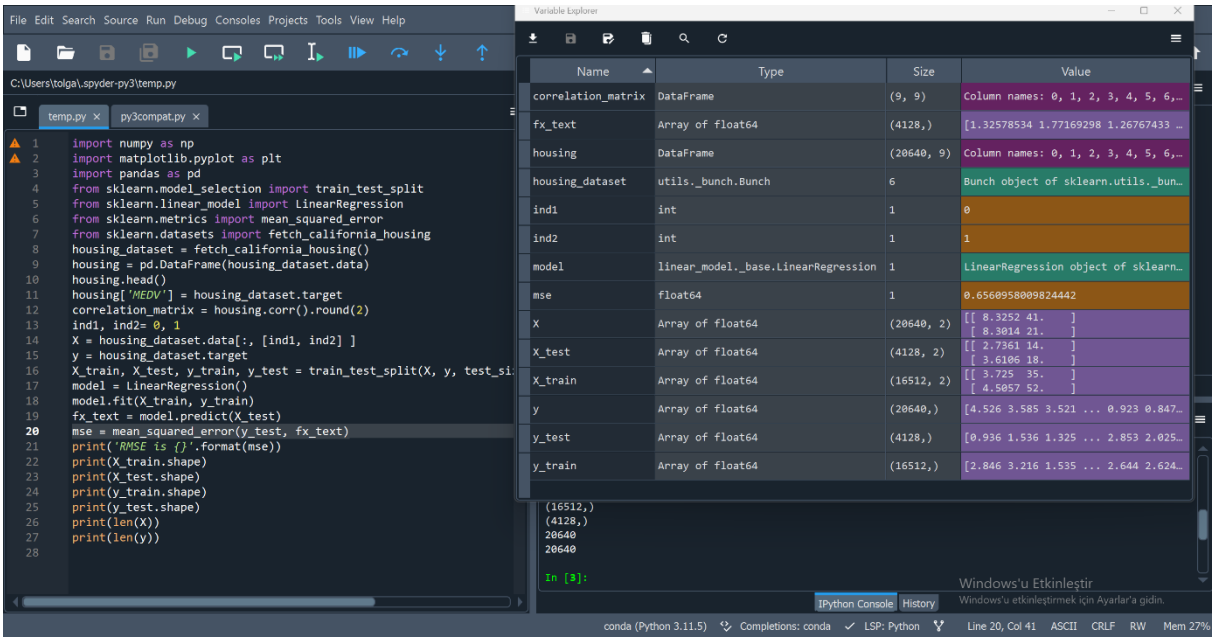
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3\temp.py
temp.py
py3compat.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
housing['MEDV'] = housing_dataset.target
correlation_matrix = housing.corr().round(2)
ind1, ind2= 0, 1
X = housing_dataset.data[:, [ind1, ind2]]
y = housing_dataset.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_si
model = LinearRegression()
model.fit(X_train, y_train)
fx_text = model.predict(X_test)
mse = mean_squared_error(y_test, fx_text)
print('RMSE is {}'.format(mse))
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
print(len(X))
print(len(y))
Variable Explorer
Name
Type
Size
Value
correlation_matrix
DataFrame
(9, 9)
Column names: 0, 1, 2, 3, 4, 5, 6,...
fx_text
Array of float64
(4128,)
[1.32578534 1.77169298 1.26767433 ...
housing
DataFrame
(20640, 9)
Column names: 0, 1, 2, 3, 4, 5, 6,...
housing_dataset
utils._bunch.Bunch
6
Bunch object of sklearn.utils._bun...
ind1
int
1
0
ind2
int
1
1
model
linear_model._base.LinearRegression
1
LinearRegression object of sklearn...
mse
float64
1
0.6560958009824442
X
Array of float64
(20640, 2)
[[8.3252 41.]
[8.3014 21.]
X_test
Array of float64
(4128, 2)
[[2.7361 14.]
[3.6106 18.]
X_train
Array of float64
(16512, 2)
[[3.725 35.]
[4.5057 52.]
y
Array of float64
(20640,)
[4.526 3.585 3.521 ... 0.923 0.847...
y_test
Array of float64
(4128,)
[0.936 1.536 1.325 ... 2.853 2.025...
y_train
Array of float64
(16512,)
[2.846 3.216 1.535 ... 2.644 2.624...
(16512,)
(4128,)
20640
20640
In [3]:
IPython Console
History
Windows'u Etkinleştir
Windows'u etkinleştirmek için Ayarlar'a gidin.
conda (Python 3.11.5)
Completions: conda
LSP: Python
Line 20, Col 41
ASCII
CRLF
RW
Mem 27%

1.9)

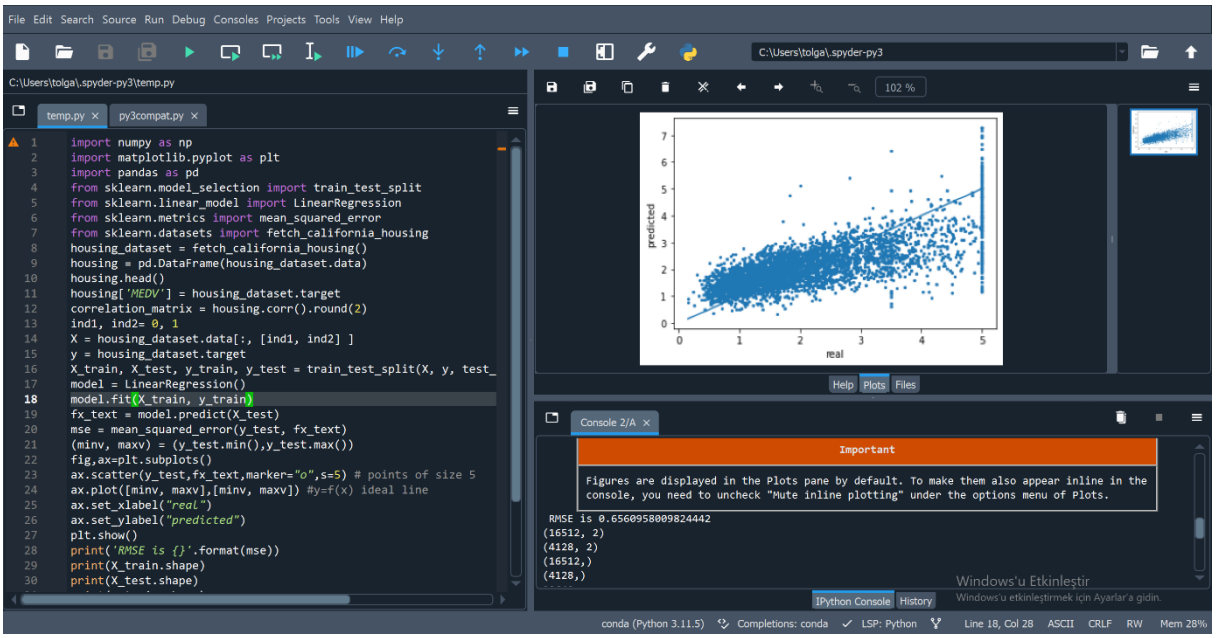
File Edit Search Source Run Debug Consoles Projects Tools View Help
C:\Users\tolga\.spyder-py3
C:\Users\tolga\.spyder-py3\temp.py
temp.py
py3compat.py
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
housing_dataset = fetch_california_housing()
housing = pd.DataFrame(housing_dataset.data)
housing.head()
housing['MEDV'] = housing_dataset.target
correlation_matrix = housing.corr().round(2)
ind1, ind2= 0, 1
X = housing_dataset.data[:, [ind1, ind2]]
y = housing_dataset.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_
model = LinearRegression()
model.fit(X_train, y_train)
fx_text = model.predict(X_test)
mse = mean_squared_error(y_test, fx_text)
(minv, maxv) = (y_test.min(),y_test.max())
fig,ax=plt.subplots()
ax.scatter(y_test,fx_text,marker="o",s=5) # points of size 5
ax.plot([minv, maxv],[minv, maxv]) #y=f(x) ideal line
ax.set_xlabel("real")
ax.set_ylabel("predicted")
plt.show()
print('RMSE is {}'.format(mse))
print(X_train.shape)
print(X_test.shape)
102 %
predicted
real
Help
Plots
Files
Console 2/A
Important
Figures are displayed in the Plots pane by default. To make them also appear inline in the console, you need to uncheck "Mute inline plotting" under the options menu of Plots.
RMSE is 0.6560958009824442
(16512, 2)
(4128, 2)
(16512,)
(4128,)
IPython Console
History
Windows'u Etkinleştir
Windows'u etkinleştirmek için Ayarlar'a gidin.
conda (Python 3.11.5)
Completions: conda
LSP: Python
Line 18, Col 28
ASCII
CRLF
RW
Mem 28%

2.1)

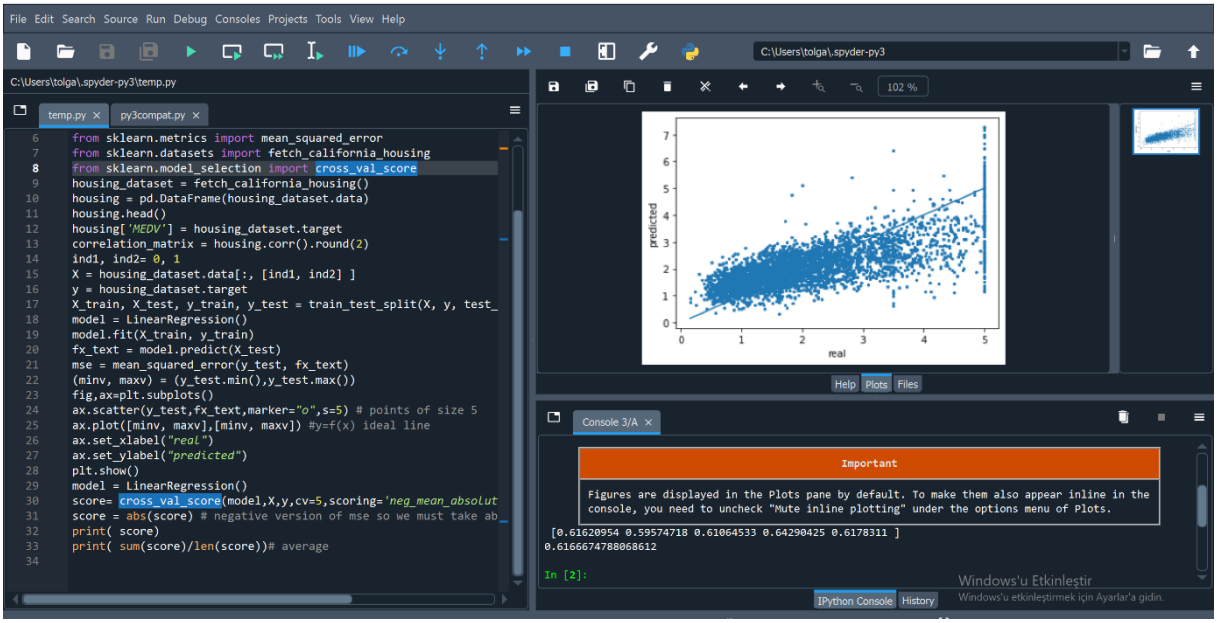

If the average cross-validation MSE is close to the validation set MSE, it suggests that the model is generalizing well to different subsets of the data.

If the average cross-validation MSE is significantly lower than the validation set MSE, it could indicate that the model is overfitting to the training set, and its performance doesn't generalize well to

2.2)